

Supervised Learning with Black Box Loss Functions

David S. Rosenberg

NYU: CDS

March 31, 2021

Contents

- 1 Black box and non-differentiable losses
- 2 Randomization and policy gradient
- 3 Policy gradient for sequence prediction
- 4 Image to Sequence

Black box and non-differentiable losses

Supervised learning

- Input space \mathcal{X} ; Label space \mathcal{Y}
- Hypothesis space of functions $x \mapsto f_{\theta}(x) \in \mathcal{Y}$
- For $(X, Y) \sim P$, loss function $\ell(f_{\theta}(X), Y)$ tells us how we did.
- For gradient-based learning methods, we compute

$$\nabla_{\theta} \ell(f_{\theta}(X), Y) = \frac{\partial \ell}{\partial f_{\theta}(X)} \ell(f_{\theta}(X), Y) \nabla_{\theta} f_{\theta}(X)$$

and iteratively update θ as $\theta \leftarrow \theta - \eta \nabla_{\theta} \ell(f_{\theta}(X), Y)$

What about non-differentiable or black-box losses?

- For gradient-based learning methods, we compute

$$\nabla_{\theta} \ell(f_{\theta}(X), Y) = \frac{\partial \ell}{\partial f_{\theta}(X)} \ell(f_{\theta}(X), Y) \nabla_{\theta} f_{\theta}(X)$$

- What if the loss function ℓ is not differentiable? not continuous?
- What if the loss function ℓ is a black box?
- Examples:
 - 0/1 loss for classification
 - BLEU score – used in machine translation for scoring a candidate sentence against a reference set of translations
 - ROUGE metrics – used in automatic summarization for scoring a summaries against a reference set of summaries
 - a human annotator evaluating predictions made by a model

The essence of the issue

- Gradient methods are iterative and “local”.
- Suppose our current model is f_θ .
- By examining $\ell(f_\theta(X), Y)$ in a small neighborhood of θ , we can figure out how to make a small change to θ to improve performance.
- We can do this conveniently with the gradient.
- But with discrete label spaces
 - $\theta \mapsto \ell(f_\theta(X), Y)$ will be piecewise constant, and so.
 - we get no information about how to change θ by looking in a small neighborhood of θ .
- So the situation can be much worse than just “non-differentiable”.

Randomization and policy gradient

Using randomness to smooth our objective

- Let's move from deterministic predictions $f_{\theta}(x) \in \mathcal{Y}$
- to randomized actions $A \in \mathcal{Y}$ drawn from $\pi_{\theta}(a | x)$.
- Let's consider the expected loss as our performance measure:

$$\mathbb{E}_{A \sim \pi_{\theta}(a | X)} \ell(A, Y).$$

- We want a conditional probability model on actions that gets small expected loss.
- How has this helped?

Expected loss for randomized actions

- Expected loss is

$$\mathbb{E}_{A \sim \pi_\theta(a|X)} \ell(A, Y) = \sum_{a \in \mathcal{A}} \pi_\theta(a|X) \ell(a, Y).$$

- Before, we were evaluating the loss at $\ell(f_\theta(X), Y)$, where
 - $f_\theta(X)$ may change discontinuously as a function of θ
 - Can't figure out what direction to move θ from a local neighborhood of θ
- For expected loss, we're always evaluating the loss on all possible actions
 - As θ varies, we're changing the relative weights on losses from each action
- Expected loss changes smoothly as θ changes (so long as $\pi_\theta(a|X)$ changes smoothly).

Gradient of expected loss

- Gradient of expected loss:

$$\begin{aligned}\nabla_{\theta} \left[\mathbb{E}_{A \sim \pi_{\theta}(Y|X)} \ell(A, Y) \right] &= \nabla_{\theta} \left[\sum_{a \in \mathcal{A}} \pi_{\theta}(a|X) \ell(a, Y) \right] \\ &= \sum_{a \in \mathcal{A}} \ell(a, Y) \nabla_{\theta} \pi_{\theta}(a|X)\end{aligned}$$

- We can compute the gradient of expected loss so long as the CPM is differentiable.
- We don't need to differentiate w.r.t. the loss.
- What if the action space \mathcal{A} is too large to sum over?

Clever trick again

- We have

$$\begin{aligned}\nabla_{\theta} \left[\mathbb{E}_{A \sim \pi_{\theta}(a|X)} \ell(A, Y) \right] &= \sum_{a \in \mathcal{A}} \ell(a, Y) \nabla_{\theta} \pi_{\theta}(a | X) \\ &= \sum_{a \in \mathcal{A}} \ell(a, Y) \pi_{\theta}(a | X) \nabla_{\theta} \log \pi_{\theta}(a | X) \\ &= \mathbb{E}_{A \sim \pi_{\theta}(a|X)} \ell(A, Y) \nabla_{\theta} \log \pi_{\theta}(A | X)\end{aligned}$$

- Now we can use Monte Carlo to estimate this.
- If we have samples A_1, \dots, A_n from $\pi_{\theta}(a | X)$, our action-generating distribution, then

$$\frac{1}{n} \sum_{i=1}^n \ell(A_i, Y) \nabla_{\theta} \log \pi_{\theta}(A_i | X)$$

- is an unbiased estimate of $\nabla_{\theta} \left[\mathbb{E}_{A \sim \pi_{\theta}(a|X)} \ell(A, Y) \right]$.
- Look familiar?

Comparison to policy gradient for contextual bandits

- We have rederived policy gradient for contextual bandits.
- Are there any differences?
- We get the ground truth label Y .
- If we also have access to the loss function (at least as a black box),
 - we can feed in lots of possible actions for the given X .
- This will give us a much better estimate of $\nabla_{\theta} [\mathbb{E}_{A \sim \pi_{\theta}(a|X)} \ell(A, Y)]$.
- Compared to policy gradient, where we observe the reward for only a single action.
- We can also directly apply policy gradient, using just a single action to estimate the gradient.

Comparison to maximum likelihood

- We're using conditional probability models.
- And we observe ground truth labels.
- Why not just fit the CPMs with maximum likelihood?
- That's typical, but it ignores our loss function!
- Perhaps some errors are worse than others.
- Maximum likelihood focuses on putting as much weight as possible on the correct label.
- But maybe the loss function is fairly indifferent between multiple labels.
 - e.g. There may be multiple translations that are essentially equivalent, but our "label" is just one of them.
 - Maximum likelihood tries to put all the weight on the "correct" one.
 - Maximizing w.r.t. a loss function that's indifferent to these variations won't do this.
- On the negative side, policy gradient estimates are very high variance, and
 - optimization is very slow.

Policy gradient for sequence prediction

Application: Sequence-to-Sequence Models

- Consider machine translation.
- e.g. Conditioned on sentence in English, produce a distribution on sentences in French.
- Model is $\pi_{\theta}(a | x)$, where x is an English sentence and a is a French sentence.
- Typically trained on ground truth pairs (X_i, Y_i) using maximum likelihood:

$$\theta^* = \arg \max_{\theta} \prod_{i=1}^n \pi_{\theta}(Y_i | X_i).$$

- Seems reasonable...
- But how do we actually measure performance for machine translation?

Application: Sequence-to-Sequence Models

- Suppose we are assessing performance of our MT model on a test set.
- We get input X .
- We use our current model $\pi_{\theta}(\cdot | X)$ and produce a sequence A .
 - (e.g. by sampling or beam search or something)
- Suppose A is a perfect translation of X , but it's different from the ground truth Y .
- We'd like to give credit for this translation.
- I don't think there's a great way to do this in an automated way.

But there is BLEU score

- A frequent measure of translation quality is BLEU score.
- Let's not discuss the details of BLEU score.
- For our purposes, sufficient to know that
 - BLEU takes a proposed translation and a ground truth and gives a numerical score
- BLEU score is computed by an algorithm and is **not differentiable**.
- Perhaps it would make sense to train a model to optimize directly for BLEU score?

- Sequence models are frequently **autoregressive**.
- We condition on previously predicted tokens to predict the next token in a sequence.
- During MLE training, we're always conditioning on the gold labels in Y .
- During test, we're conditioning on our own predicted labels.
- **Our model never trains using its own predictions as input.**

- This is a known issue with maximum likelihood training of sequence models.
- There is a family of approaches called “learning to search” that address this issue.
- e.g. SEARN, DAgger, AggreVaTe, LOLS, etc.
- We can use policy gradient as well.
- Sampling from $\pi_{\theta}(\cdot | X)$ doesn't involve the ground truth sequence Y at all.
- We don't even have the ground truth label to use, except as part of the reward function.

Usually we pre-train with maximum likelihood



- Suppose we want to train a sequence-to-sequence model
 - with BLEU score as reward.
- Policy gradient is sufficient for this task.
- In practice, we usually pretrain our model with maximum likelihood.
 - It's much faster than policy gradient.
- Then switch to policy gradient to optimize to our particular loss/reward.

Self-Critical Baseline

- When we have access
 - to the loss function $\ell(a, y)$ and
 - the ground truth labels Y_1, Y_2, \dots ,
- there's another clever way to set a baseline:
 - Find (or approximate) the action that is optimal under our policy: $A_t^* \approx \arg \max_a \pi_{\theta_t}(a|X_t)$
 - Use the loss $\ell(A_t^*, Y_t)$ as a baseline.
- If the current action performs better than the action our policy says is best, then we should make the current action more likely.
- If it performs worse than what the policy says is best, let's make it less likely.

Image to Sequence

Image to Sequence Problems¹



```
<object>
  <supercategory>C-1</supercategory>
  <category>CS-3</category>
  <x-coordinate>120</x-coordinate>
  <y-coordinate>240</y-coordinate>
  <depth>1</depth>
  <flip>0</flip>
</object>
<object>....
```

```
<object>
  <category>Rectangle</category>
  <x1-coordinate>7</x1-coordinate>
  <y1-coordinate>1</y1-coordinate>
  <x2-coordinate>11</x2-coordinate>
  <y2-coordinate>16</y2-coordinate>
</object>
<object>....
```

¹Results and images for this section are taken from [PRMB21]

How to Evaluate?

- One obvious idea is to re-render and check for an exact match.
- This is a very challenging metric.
- We only get positive feedback when we get the image exactly correct.
- Will take a **very long time** to learn this way (at least starting from scratch).
- Doesn't work for hand-drawn shapes

- Two specifications can be very different, yet render to very similar things.
 - e.g. by reordering objects
- Two images may look very different (e.g. at the pixel level), but have similar specifications
 - e.g. by changing a color
- We can evaluate performance in **image space** and in **specification space**.

Image Space Measure

- We can measure performance in image space with

$$d_{img} = \|I - \Psi(I^R)\|_2^2,$$

where I is the original image vector and I^R is the rendering of the predicted image.

- For the noisy shapes dataset, Ψ is a Gaussian blurring function.
- For the abstract scene dataset, Ψ is identity function.
- Why isn't this differentiable?
- Computing I^R uses a graphics renderer...
- (There are differentiable renderers now... but that's another story.)

Specification Space Measure: IOU Reward

- Our specifications break down into “objects”.
- We can look for exact matches between prediction and ground truth at the object level.
- For numeric attributes, we divide range into 20 bins of equal size
 - consider it a match if the bin is correct
- Can summarize matches with precision, recall, F1, etc.
- A common summary in this scenario is **intersection-over-union** (IOU)....


Intersection over Union

- Let $\{o_i\}_{i=1}^m$ and $\{o_j^*\}_{j=1}^n$ represent the objects in predicted and ground-truth specifications, respectively.
- Then the IOU reward is defined as follows:

$$r_{iou} = \frac{\text{count}(\{o_i\}_{i=1}^m \cap \{o_j^*\}_{j=1}^n)}{\text{count}(\{o_i\}_{i=1}^m \cup \{o_j^*\}_{j=1}^n)}$$

- Roughly speaking, IOU gives credit for predicting objects that exactly match objects in the ground truth
- Penalizes both for predicting objects that do not match ground truth objects and for failing to predict objects that are part of the ground truth.

Model: ResNet to Transformer Decoder



Results: cross-entropy Loss (i.e. maximum likelihood)

Model	Recons. IOU Error	
Cross-Entropy Loss		
Image2LSTM+atten.	15.70	32.06
Image2Transformer	10.92	58.54

- reconstruction error corresponds to the image distance
- average error across test set

Results: policy gradient

Model	Recons. IOU Error	
Cross-Entropy Loss		
Image2LSTM+atten.	15.70	32.06
Image2Transformer	10.92	58.54
Image2Transformer with Reinforce Loss		
IOU Reward	10.50	61.29
Recons. Reward	9.99	62.44
IOU + Recons.	10.04	62.45

References

- The idea of the self-critical baseline comes from [RMM⁺17].

- [PRMB21] Ramakanth Pasunuru, David S. Rosenberg, Gideon Mann, and Mohit Bansal, *Dual reinforcement-based specification generation for image de-rendering*, Proceedings of the Workshop on Scientific Document Understanding co-located with 35th AAAI Conference on Artificial Intelligence, SDU@AAAI 2021, Virtual Event, February 9, 2021 (Amir Pouran Ben Veyseh, Franck Dernoncourt, Thien Huu Nguyen, Walter Chang, and Leo Anthony Celi, eds.), CEUR Workshop Proceedings, vol. 2831, CEUR-WS.org, 2021.
- [RMM⁺17] Steven J. Rennie, Etienne Marcheret, Youssef Mroueh, Jerret Ross, and Vaibhava Goel, *Self-critical sequence training for image captioning*, 2017 IEEE Conference on Computer Vision and Pattern Recognition (CVPR), 7 2017, p. nil.